

# Project Proposal: Distributed Sharded Key-Value Store

Group 41: Ruichen Zhang, Roger Wu, Kiran Pyles

April 5, 2025

## 1. Project Overview

The objective of this project is to design and implement a distributed key-value store that leverages sharding to achieve scalability, high performance, and fault tolerance. The system will store data as key-value pairs, distributing the dataset across multiple nodes (shards) using techniques such as consistent hashing. By integrating core distributed system concepts, the project aims to provide a practical demonstration of horizontal scaling, load balancing, and resilience against node failures.

## 2. Goal and Scope

### Goal

Develop a fully functional distributed key-value store capable of handling high throughput for read and write operations, while dynamically managing data distribution as nodes are added or removed.

### Scope

- **Data Distribution:** Implement sharding using consistent hashing to evenly distribute keys.
- **Fault Tolerance:** Ensure system resilience by replicating data across shards.
- **Scalability:** Demonstrate horizontal scaling by adding nodes with minimal data rebalancing.
- **Performance:** Optimize for low latency in key lookup and update operations.
- **Consistency:** Evaluate tradeoffs between eventual consistency and stronger consistency models under various network conditions.

## 3. System Architecture

The architecture consists of several key components:

1. **Client Interface:** Clients interact with the system using a simple API to perform put, get, and delete operations on keys.
2. **Coordinator/Router:** A routing layer maps incoming keys to the appropriate shard using a consistent hashing algorithm. This layer redirects client requests to the correct nodes.
3. **Shard Nodes:** Each shard is a server instance responsible for managing a subset of key-value pairs. Nodes communicate to replicate data for fault tolerance.
4. **Replication Mechanism:** Data will be replicated across multiple nodes. A replication strategy will ensure data availability in case of node failure.
5. **Rebalancing Service:** A background process will manage data redistribution when nodes are added or removed, minimizing data movement and system downtime.

## 4. Implementation Plan

### 1. Design Phase:

- Define API specifications for client-server interactions.
- Design the consistent hashing mechanism and replication strategy.
- Create system architecture diagrams and flowcharts for key operations.

### 2. Development Phase:

- Implement the client interface and coordinator/router.
- Develop shard node functionality for storage, replication, and recovery.
- Integrate the rebalancing service to support dynamic cluster changes.
- Write unit and integration tests to ensure reliability.

### 3. Testing Phase:

- Simulate network conditions such as node failures and high load.
- Benchmark performance for read/write operations and rebalancing efficiency.
- Adjust the replication and consistency settings based on test outcomes.

### 4. Documentation & Deployment:

- Provide comprehensive documentation covering system design, API usage, and deployment instructions.
- Deploy a demo instance for live testing and presentation.

## 5. Potential Challenges and Solutions

### • Data Rebalancing Overhead:

*Challenge:* Adding or removing nodes may trigger significant data movement.

*Solution:* Implement incremental rebalancing using virtual nodes to limit disruption.

### • Network Partitions:

*Challenge:* Ensuring consistency during network splits.

*Solution:* Use configurable consistency models (e.g., eventual vs. strong consistency) and implement quorum-based reads/writes.

### • Replication Lag:

*Challenge:* Delay in propagating changes across nodes could impact read consistency.

*Solution:* Optimize replication protocols and monitor lag, integrating conflict resolution mechanisms if necessary.

## 6. Final Deliverables

- **Source Code:** Fully documented codebase for the distributed key-value store.
- **Design Documentation:** Detailed system architecture, API specifications, and design rationale.
- **Test Suite:** Automated tests covering unit, integration, and performance scenarios.
- **Demo Deployment:** A live instance or simulation demonstrating system scalability, performance, and fault tolerance.
- **Final Report/Presentation:** A comprehensive report summarizing design decisions, evaluation results, encountered challenges, and future improvements.